

FRED TALKS

23rd April 2026

CONTENTS

Thoughts and Feelings around Claude Design

SAMHENRI.GOLD · 1230 WORDS

Apr 22, 2026Economic ResearchAnnouncing the Anthropic
Economic Index Survey

ANTHROPIC RESEARCH · 588 WORDS

Apr 22, 2026Economic ResearchWhat 81,000 people told us about
the economics of AI

ANTHROPIC RESEARCH · 2254 WORDS

The Slide

RANDS · 775 WORDS

The commodification of travel

HERMAN (HIDDEN) · 582 WORDS

Thoughts and Feelings around Claude Design

SAMHENRI.GOLD · 19 APR 2026 · [SOURCE](#)



I tried [Claude Design](#) yesterday and I have a theory for how this whole thing shakes out.

As product teams scaled and design needed to justify itself inside engineering orgs, it was pushed toward systematization — and Figma invented its own primitives to make that work: components, styles, variables, props, and so on. Some concepts are borrowed from programming, some aren't, and the whole thing doesn't neatly map onto anything. Guidance evolves, migrations pile up, and if you want to automate any of it you're stuck with a handful of shoddy plugins. The beast is hairy enough that entire design roles now specialize in wrangling the system itself.

There's always been a tense push-pull between Figma and code over what the source of truth should be. Figma won over Sketch partially by staking its claim there — their tooling would be canonical.

That victory had a hidden cost. By nature of having a locked-down, largely-undocumented format that's painful to work with programmatically, Figma accidentally excluded themselves from the training data that would have made them relevant in the agentic era. LLMs were trained on code, not Figma primitives, so models never learned them. As code becomes easier for designers to write and agents keep improving, the source of truth will naturally migrate back to code. And all the baroque infrastructure Figma had to introduce over the past decade will look nuts by comparison. Why fuss around in a lossy approximation of the thing when you can work directly in the medium where it will actually live? If we want to make pottery, why are we painting watercolors of the pot instead of just throwing the clay?

I received a gift many years ago which, while being well-intentioned, typifies this idea: a scratch-off map of the world. Each time you visit a country, you can scratch off the metallic coating and the country is now *done*, according to the map. The work trip I took to São Paulo a decade ago? Brazil: done. Bus tour through Europe? Germany? Check. France? Check. Spain? You get the idea.

This kind of mentality is typified in the question I've heard asked many times: "How many countries have you been to?" This is often followed by a debate on whether layovers count towards your tally if you don't leave the airport, as if stepping beyond the airport boundaries bestows *doneness*.

Like many things, I blame social media. It's changed travel from an exploration to social status signalling. I started thinking about this a few years ago while visiting some waterfalls in Indonesia. I love a good frolic in a waterfall, but all of them were just lines of people waiting to take their photo under the falls, and then they'd better get out of the way for the next photo-goers. No frolicking allowed! People need to *do* these waterfalls!

I spent this morning in a beautiful garden outside of Kyoto, which exemplifies the cultural ideals of appreciating nature and meditating on the beauty that surrounds us. It was lovely during the early morning, but then the rest of the world showed up and all they wanted to do was take photos and move on to the next spot to *do*. There was one moment, in perhaps one of the most heart-wrenchingly beautiful places I've ever visited, where I was surrounded by about 20 other people, all of them either in the process of taking a photo, or looking at what they had just taken on their phones. No one was looking at the amazing stuff they were *doing*!

That isn't to say taking photos is bad. They're a great way to share an experience with others and save a memory of a time and place—but I think the threshold of what is *enough* has been crossed in the age of Instagram where images and video are socially valuable. Now beautiful places are commodified. And I don't know if we'll ever go back.

I appreciate that many places in Japan limit photos and videos, such as on trains or in gyms, for the sake of not annoying those around you. Perhaps once sunglasses cameras take off and people can record their entire lives they can finally experience where they are, instead of trying to capture it perfectly for later.

All that being said, I don't want to gate-keep. If this is the form of travel that makes people happy, then they should do it to their heart's content. Similarly to how some people collect Magic cards while never playing the game, sometimes the fun is in the collection itself. But perhaps look up from your phone once in a while. The world is prettier in full resolution.

So, what feedback had I been attempting and failing to give this mysterious former manager prior to this bearing of my soul? Correct. Delegation. The single biggest challenge for new managers — giving up the responsibility for the product... for the building. Learning how to give accountability for projects of significance to the team. It's an essential set of complex skills involving trust, communication, and, most importantly, judgment. Failure to understand delegation is failing to be a leader. Senior or not.

My thesis is why this skill is hard to learn; the reason they don't want to hear this feedback is that it contradicts the valuable core engineering skills that got them the role in the first place. The Slide is you gently sliding up right next to that discomfort, that contradiction, and not accusing, not lecturing, just telling the story of that time you learned the thing.

Not a Guarantee

Why won't they listen? What is it about this particular habit or behavior that has this capable, smart, and reasonable human ignore the advice of a seasoned, well-informed, and trusted leader?

The answer is usually fear. The variants of fear that apply here are as numerous as the situations, but fear is fear. They have an inner monologue about this topic, "I will be less if I do this. I will have failed if I don't achieve. I should have known. They will finally know I am a fraud."


You will never diagnose the fear, but slide up next to them and tell them about the time you were scared, too.

The commodification of travel


HERMAN (HIDDEN) · 22 APR 2026 · [SOURCE](#)

I've noticed that travel has become, of late, an act of collecting places. I've literally heard people referring to visiting a place as *doing* that place, as in "Have you done Japan?", assuming that one can *do* an entire country, and once that country is *done* it remains as such. As if a place is a product to be consumed and checked off the list. Why bother returning to a place if you've already *done* it?


At work, we've spent quite a bit of time back-porting design changes made directly in code back to Figma and it is not fun. I can't share that file, but for a fair comparison, this is Figma's own design system file for their product. I have to assume it was built by the most competent design system team you can find. And yet...




The Figma variables panel showing 946 color variables organized into nested groups like "bg/desktopBackgrounded," with a single selected variable revealing eight mode-specific values: Light, Dark, FigJam-Light, FigJam-Dark, DevMode-Light, DevMode-Dark, Slides-Light, and Slides-Dark.



A modal footer component open to its variant property editor, showing 12 variants with a dropdown full of values like "DS Library Swap," "QA Plugin," "Growth Stepper," and "Sharing Actions." The right panel lists eight props like "Border," "Second CTA," and "Helper Text"



The effect styles panel for a slider component, showing a style named "light/elevation-300-tooltip." Expanding it reveals its entire definition: a 0.5px drop shadow at 30% black. It has its own named style because that's the only way to document what CSS variable it corresponds to.



A combo input component with 16 variants. Its children in the layers panel are named things like "Default, Default, Close Button=False" and "Default, Focused, Close Button=True"

These are Figma's own files. Built by their own team. This is the gold standard.

Imagine debugging a color that looks wrong. You check the component. The component uses a variable. The variable is aliased to another variable. That variable references a mode. The mode is overridden at the instance level. The instance lives inside a nested component with a library swap applied. At this point, you're either considering picking up code or moving to the countryside and becoming a sheep farmer because one more minute of this will make you lose your goddamn mind.

So as the source of truth shifts back to code, Figma is left in an odd spot: holding a largely manual, pre-agentic system that nobody in their right mind would design from scratch today.

I think design tooling forks into two distinct shapes from here — and there's almost a clock resetting between Figma and every other tool competing to answer the same question they answered in 2016: who can help me, a designer, get my ideas out fastest?

Not a Trick

I have a move. It's not a guarantee, but if you've tried the obvious approaches, if you've tried straight talk, and if you've made no progress, I offer The Slide.

There are prerequisites for The Slide:

1. You've tried a couple of different approaches to giving this feedback. They've heard it, but they have not acted.
2. You have high confidence that if they actually absorbed the feedback and acted on it, they'd attempt to change. Somewhere in the back of their head, past the denial, you know they'll get it.
3. You've had to learn the same lesson in your professional career.

I can not tell you when to deploy The Slide, the opportunity will present itself when the person who needs the feedback, once again, complains or otherwise comments on the by-product consequence of their negligence. Yes, it's infuriating because if they just listened to you, they'd have a stronger set of tools to tackle the problem, but bury that and Slide 'em:

Them: "Yeah, and isn't just the endless meetings, it's the fact that I don't have anyone on the team who can do the meetings. Francis is deep in backend debugging, Jake isn't ready to run that meeting, and Jason, well, Jason doesn't know what he doesn't know."

Sweet, sweet irony. Take a deep breath and Slide:

"Back at Pinterest, we didn't have a CTO. I was the VP of Engineering, and had this loose collection of very bright senior engineers who wanted to help. Problem was... me. Whenever a CTO-class problem came up, I'd try to be the CTO, which meant I wasn't being the VP of Engineering. After a bit, I was doing poorly at both jobs."

Them: "Poorly, how? What'd you do?"

"It's mostly the sense of 'Is there enough time in the day?' If the answer to that question is ever 'no', then I'm doing poorly. Doesn't hurt when others point that out, too. After a few months of barely treading water, I gathered together the senior engineer leaders, and we built a small council. When CTO problems arrived, we gave it one of them. They drove, but they relied on the other members and me to get it done."

productive, 5 = substantially more productive, 6 = much more productive, to 7 = transformatively more productive—AI has fundamentally changed what or how much they can produce.

4. Removing these “solopreneurs” still leaves management tied with computer and math occupations for showing the highest productivity benefit.
5. A major caveat, however, is that this survey went out to people with personal Claude accounts. A more representative picture would also include enterprise users, who may be more likely to say the value accrues to their employers.

The Slide

RANDS · 22 APR 2026 · [SOURCE](#)

They don’t want to hear it.

They don’t.

It is legit good feedback. It is substantive, and it is 100% true, but they do not want to hear it.

How do you know this? You’ve provided the feedback several times. Three different variants of the critical feedback, and each time, the response is one of these:

- “I know, I know. I’m working on it.” *They are not.*
- An immediate, abrupt topic change. *What?*
- (My favorite) Twisting the feedback and making it about you. *Impressive.*

This is not a bad employee, this is not a person incapable of changing, and this is most certainly not an adversarial situation.

They don’t want to hear it.

Spoiler: it’s not Figma Make. Figma Make feels like it primarily benefits people who have already drunk the Kool-Aid — it reads from Figma styles, component libraries, and proprietary props (or, as I like to call them, Prop Props), and it’s the only tool in this new landscape still pretending the design file is canonical. It’s the tool for people who want to (or have no choice but to) stay inside the system.

Claude Design is the first of those two tools, and takes the opposite bet. There’s an Arts and Crafts principle called “[truth to materials](#)” — the idea that a thing should be honest about what it is and how it’s made, rather than masquerading as something else. Figma ended up being the opposite of this: a set of extremely rigid schemas with a free-form “just vibes, man” costume over the top. Like a Type-A personality physically incapable of relaxing, forced to perform chill while internally screaming that your frames aren’t nested and your tokens are detached and nothing is on the grid. Claude Design, for all its roughness, is at least honest about what it is: HTML and JS all the way down.



A Gustav Stickley lamp table, circa 1902. The joinery is exposed, not hidden. The wood is the wood.

And it has a massive structural advantage: its sibling is Claude Code. Eventually, I can see Claude Design just dumping things directly into Claude Code and vice versa. Claude Design’s onboarding already lets you import your repos. The feedback loop between design and implementation — which has been a source of friction since the beginning of time — becomes a single conversation.

The other tool that emerges from this moment will have no expectation of code at all. It’ll be a pure exploration environment — somewhere to drop rectangles, stack layer styles, fuss with blend modes and gradients, and go completely nuts, unconstrained by systems or prompting conventions. Maybe it’s an iPad app with Pencil support where you just quickly sketch a bunch of rectangles. [37signals could do something really funny right now](#). Or maybe it goes in the opposite direction — something more like Photoshop that goes all-in on high-fidelity compositing and lets our imaginations run wild, now that we’re no longer beholden to the ceiling of what you can do with CSS effects. Doesn’t it seem *kinda weird* how for 90% of its life, Figma’s only layer effect was a drop shadow or a blur?

Figma’s Sketch moment is rapidly approaching. And if you said that sentence to a Victorian child, they would probably have a stroke.

Post Script

The following are messages meant only for the teams behind Sketch and Figma. If neither apply to you, you can skedaddle.

To Figma: I can see a world where this post does numbers in the Figma internal Slack. If that’s the case and you’re reading this from Figma: this wouldn’t have happened if you hired me last year when I was interviewing. Your loss, big dawg.

To Sketch: GET YOUR HEADS OUTTA YOUR ASSES AND GIVE EM HELL. ADD PARTICLE EFFECTS. ADD DEBOSSING EFFECTS. MESH TRANSFORMS. FUCK IT, ADD METAL SHADERS. GO NUTS. STOP COASTING OFF OF BEING MAC NATIVE. QUIT DRINKING COCOA AND GET THIRSTY FOR BLOOD.


To mom: Sorry for cursing.

Post-Post Script

[@jonnyburch on Twitter](#) shared a link to [their blog post](#) with similar thoughts, it’s quite good if you wanna go deeper.

Apr 22, 2026Economic ResearchAnnouncing the Anthropic Economic Index Survey

ANTHROPIC RESEARCH · 22 APR 2026 · [SOURCE](#)



users weren’t asked directly about many of the derived variables here, so our inferences on occupation, career stage, and other variables from contextual clues could be wrong. Relatedly, because the survey is open-ended, our measures are based on what respondents happen to mention; these findings should be confirmed in structured surveys that ask about these topics directly.

Still, the interviews surface real insights about people’s feelings around the economics of AI, showing how qualitative data can surface quantitative hypotheses. And the large share of economic-related concerns is a strong signal in itself.

See the final section of the [linked PDF](#).

We thank the 80,508 Claude users who shared their stories.

Maxim Massenkoff led the analysis and wrote the blog post. Saffron Huang led the interview project and provided guidance throughout.

Zoe Hitzig and Eva Lyubich provided critical feedback and methodological guidance. Keir Bradwell and Rebecca Hiscott gave editorial support. Hanah Ho and Kim Withee contributed to design. Grace Yun, AJ Alt, and Thomas Millar implemented Anthropic Interviewer within Claude.ai. Chelsea Larsson, Jane Leibrock, and Matt Gallivan contributed to survey and experience design. Theodore Sumers contributed to the data processing and clustering infrastructure. Peter McCrory, Deep Ganguli, and Jack Clark provided critical feedback, direction and organizational support.

Additionally, we thank Miriam Chaum, Ankur Rathi, Santi Ruiz, and David Saunders for their discussion, feedback, and support.

1. We inferred people’s occupations using the first question in the survey (“What’s the last thing you used an AI chatbot for?”) or indications given in other responses.
2. This came from various indications in the written responses. For example, several users mentioned using Claude for homework, which put them in the early-career group. And many referred to running their own businesses and being involved in hiring decisions, which put them in the senior group.
3. The scale is not centered because most people say positive things about productivity, yielding almost entirely 6s and 7s on the original Likert scale. The scale we use here ran from 1 = less productive, 2 = no change, 3 = slightly more productive, 4 = moderately more

We found that the relationship between speedup and perceived job threat is U-shaped (see Figure 6). The leftmost bar shows respondents who reported that AI slowed them down. These respondents were more likely to indicate that AI posed a significant threat to their livelihoods. For example, some creative workers, like fine artists and writers, found AI too stifling and rigid to help them at their own work. At the same time, they feared the diffusion of AI into creative fields would make it harder for them to find work.




Figure 6: Job threat from AI and speedup. Percentage of respondents who said that displacement at their own job was already happening or likely in the near term, by the level of inferred speedup.

For the remaining respondents, perceived job threat increases consistently with the level of speedup implied by their answers. This makes some economic sense: if the time required to do one's tasks is shrinking quickly, there may be more uncertainty about the future viability of the role.

The Economic Index reveals what people do with AI. But another key input for understanding AI's economic impact is to hear directly from people about their experience. The responses explored here show that people's intuitions track the usage data: they worry most about AI's effect in the jobs where we observe Claude doing the most work. We also find higher levels of economic anxiety among early-career workers, which aligns with past research.

There are also signs that Claude empowers its users. People are most likely to talk about benefits flowing to themselves rather than to employers or AI companies. High-wage workers were the most enthusiastic about the productivity impacts of AI, but people with low-wage jobs and lower levels of education also reported large productivity gains. Most respondents reported that Claude enhanced their capabilities in the form of broadening the scope of their work or speeding it up. But users experiencing the largest speedups were also the most nervous about AI's job impacts.

There are key caveats to our analysis, owing to the nature of the data. First, our survey is limited to users of personal accounts on Claude.ai who chose to respond. Among other potential biases, these users could be more likely to perceive the benefits as flowing to themselves. Second, the

The [Economic Research](#) team is launching the Anthropic Economic Index Survey, a monthly survey conducted through [Anthropic Interviewer](#).

Understanding AI's economic impact requires moving beyond the quantitative data we have today. Usage and diffusion metrics tell us how AI is being deployed, and traditional labor market indicators—like employment rates, wage trends, and layoffs—track what has already happened, often with meaningful delay. Both are essential, but neither captures how people experience the changes to the economy that AI brings, nor what they expect to happen as AI capabilities advance.

While AI is poised to have large effects, there is substantial uncertainty about how AI will affect jobs, productivity, and unemployment (and on what timeline). To forecast a transition that is still unfolding, we need to hear from the people who are living through it, and we need to do so on a cadence that can identify changes as they emerge.

In the Anthropic Economic Index Survey, we aim to capture a rich new corpus of qualitative data. This effort complements a [companion report](#) that takes an economic lens to the [81,000 open-ended](#) survey responses collected through Anthropic Interviewer in December. We'll ask Claude users about if AI is changing their work today—which tasks they may be handing off, if they're seeing productivity gains, what shifts they may be observing in hiring and roles—their expectations for the future, and what they hope a well-handled transition looks like.

Collecting these data monthly will enable measurement of not just what people experience and expect, but how quickly their views shift as AI capabilities evolve. Combined with Claude usage data in a [privacy-preserving way](#), these first-hand accounts can surface change before it shows up in aggregate labor market data.

The survey launches today. Each month, we will invite a small, randomly selected group of Claude users: anyone with a personal account at least two weeks old may be invited. We will rotate the sample each month so that we hear from as broad a range of people as possible over time. If you're invited, you'll see a banner on [claude.ai](#), or get an email if you use Claude primarily on mobile. We'd love you to take part. We plan to publish insights in future Anthropic Economic Index reports and other research briefs. For more information, see the FAQ below.

How do I access the study?

From today, a small random sample of Claude users with accounts at least two weeks old will see an invitation in [claude.ai](#), the Cowork desktop app, or an email if you use Claude only on mobile. We invite a new set of Claude users each month, so if you're not selected this time, keep an eye out.

What will the study ask me?

We will use Anthropic Interviewer to ask you about your work, what changes you’ve noticed in your role or organization, how you expect that to shift over the next year, and how you hope an economy shaped by AI will look in ten years.


How will you use the data?

We will analyze the insights from this study as part of our economic research, publish our findings, and use this to improve our models and services in ways that reflect what we’ve learned. The data we collect through this study will be processed according to our [Supplemental Privacy Policy](#). We may also include de-identified responses in published findings, from users who opt in. [Learn more](#).

If you have further questions, reach out via the message icon in the lower right corner of our [Help Center](#).

Apr 22, 2026Economic ResearchWhat 81,000 people told us about the economics of AI

ANTHROPIC RESEARCH · 22 APR 2026 · [SOURCE](#)



- Our recent survey of 81,000 Claude users shows that people who work in roles that are more exposed to AI have more concerns about AI-driven job displacement. These concerns are also higher among early-career respondents.
- Those in the highest- and lowest-paid occupations report the largest productivity gains, most commonly from increases in scope (doing new tasks).




Figure 4: Where does the surplus from AI productivity go? Among respondents who named a beneficiary of their AI productivity gains, the share identifying each destination.

Respondents also shared where they experienced gains in productivity. We separate this into scope, speed, quality, and cost. For example, many people using AI for coding tasks said things like, “I’m a non tech guy but now I’m a full stack developer.” This is an expansion of scope; AI unlocks new abilities for them. In contrast, some users sped up tasks they were already doing, like the accountant who said, “I built a tool that helps me finish a financing task in 15 minutes that used to take 2 hours.” Quality gains often came from more thorough checks of code, contracts, and other paperwork. And a small share of respondents mentioned the low cost of using AI: “[I]f I hire a social media manager it’s over my budget.”

We find that the most common productivity enhancement is in scope, which was cited by 48% of users who explicitly mentioned productivity effects. 40% of users who mentioned productivity emphasized speed.




Figure 5: What kind of productivity gain do users report? Share of respondents describing each type of productivity benefit.

People’s experience with Claude might also shape their concerns about AI. To assess this, we measured the speedup reported by respondents, by extracting whether their work was now much slower (which we coded as 1), showed no change in speed (4), or had become much faster (7).




Figure 3: Inferred productivity gain by occupation. The left panel shows the mean inferred productivity benefit from AI (inferred using a Claude-powered classifier) by quartile of occupational median wage from the BLS. The right panel shows the same outcome, split by major occupational group. Error bars show 95% confidence intervals.

We look at this in more detail in the right panel of Figure 3, showing the inferred productivity gain by major occupational group. At the top are management occupations. These respondents are mostly entrepreneurs using Claude to build a business.⁴ The next highest category is computer and math, which includes software developers. The two groups exhibiting the mildest productivity improvements were workers in scientific and legal professions. Some lawyers worried about AI’s ability to follow precise instructions. For example: “I have given very specific rules about what is where, how to read a legal document, what I want it to do... but it diverges every time.”

A key question as AI diffuses through the economy is where the benefits will accrue—to workers, their managers, consumers, or corporations. Respondents indicated the recipient of these gains in about a quarter of interviews. Overall, most of these people cited benefits to themselves, through faster tasks, expanded scope, and freed-up time.⁵ But 10% of respondents who named a recipient said that employers or clients were asking for and getting more work. A smaller share mentioned benefits to AI companies, and an even smaller share said that AI would be a net negative. This depended on career stage: only 60% of early-career workers indicated that they personally benefited from AI, compared to 80% of senior professionals.

- Respondents experiencing the largest speedups from AI express higher concern about job displacement.

In order to inform the public about the economic changes we’re observing with AI, our [Economic Index](#) shares what work Claude is being asked to do, and in which jobs Claude is doing the largest share of tasks. To date, however, we’ve lacked information on how these usage patterns map onto people’s thoughts and impressions of AI.

Our recent [survey study with 81,000 Claude users](#) provides a way to connect people’s economic concerns with what we’ve quantified in Claude traffic.

The survey asked people about their visions and fears around advances in AI. Many of the thoughts that people shared touched on economic topics. We learned that many people fear job displacement—though they also feel more productive and empowered at work. In some cases, AI has enabled them to start businesses, or given them time for more important things; in others, AI feels stifling, or imposed on them by their employers.

The survey’s results provide initial evidence that [observed exposure](#) (our measure of AI displacement risk) is correlated with economic concern around AI. People in highly exposed occupations—as defined by the tasks Claude is observed performing—were more nervous about economic displacement. This is consistent with people being broadly aware of AI’s diffusion and potential impacts. We expand on our findings below.

“Well like anyone who has a white collar job these days I’m 100% concerned, pretty much 24/7 concerned about losing my job eventually to A.I.”—Software engineer.¹

One fifth of the respondents in our survey voiced concern about economic displacement. Some worried about this in the abstract: one software developer cautioned about “the possibility of AI in its current state being used to replace junior positions.” Others lamented that their jobs, or aspects of their jobs, were being automated away. One market researcher said, “In terms of improving my capability, it’s no doubt. [B]ut in the future AI may replace my work.” In some jobs, people felt it made their work harder. One software developer observed that “when AI arrived, the project managers started giving harder and harder tickets and bugs to solve.”

Throughout this report, we use Claude-powered classifiers to infer people’s attributes and sentiments from their responses. For example, many participants mention their line of work in passing or give informative details about their work life, which allows us to infer their occupation. Similarly, we quantify concerns about job loss by prompting Claude to identify and interpret direct quotes in which respondents indicate that their own role is at risk of AI-driven displacement. We give example prompts in [the Appendix](#).

Respondents' perceived threat from AI was correlated with our own measure of observed exposure, which reflects the percentage of a job's tasks for which Claude is used. A respondent was more concerned about AI when our observed exposure measure for that respondent was higher. Elementary school teachers were less worried about their own displacement than software engineers, for example, consistent with the fact that Claude usage skews toward coding tasks.

We show this in Figure 1 below. The y-axis is the percentage of respondents in a given occupation who said that AI is already replacing their role or is likely to do so soon. The x-axis is observed exposure. The plot shows that, on average, people in more exposed occupations tended to express more concern about their jobs being automated away. For every 10-percentage-point increase in exposure, perceived job threat increased by 1.3 percentage points. People in the top 25% of exposure mentioned the worry three times as often as those in the bottom 25%.




Figure 1: Perceived job threat from AI and Observed Exposure. Percentage of respondents indicating some job threat from AI vs. the Observed Exposure measure from [Massenkoff and McCrory \(2026\)](#). A respondent was coded as indicating job threat if they said their role was already being replaced or substantially reduced, or that such changes were likely in the near term (coded using Claude). The green line shows a simple linear fit.

Another important worker characteristic is career stage. In previous research, we reported tentative signs of a slowdown in the hiring of recent graduates and early-career workers in the United States. For about half of respondents in this survey, we were able to infer career stage from their answers.² We found that early-career respondents were much more likely to express concern about job displacement than senior workers.




Figure 2: Concern about economic displacement by career stage. Percentage of respondents indicating some job threat from AI, by career stage. Both fields are inferred from free-form responses using Claude-powered classifiers.

Using Claude to assess the survey responses, we rated the extent of people's self-reported productivity gains from AI on a 1–7 scale, where 1 is “less productive,” 2 is “no change,” and each subsequent level denotes a larger gain. Responses that scored 7 included testimonials like, “It used to take months to make the website I [made] in 4-5 days”; Claude gave a 5 to statements like, “What might have taken four hours was accomplished in half the time,” and a 2 to ones like, “Personally, I had AI help me fix code on a website. But it took multiple passes to get the result I was after.”³

Overall, people reported meaningful productivity gains on average. The mean productivity rating was 5.1, corresponding to “substantially more productive.” Our respondents were, of course, active Claude users who were willing to take a survey. This could make them more likely to report productivity benefits than the average user. Some 3% reported negative or neutral impacts, and 42% did not give a clear indication on productivity.

This splits somewhat across income lines. The left panel in Figure 3 shows that people in high-paying jobs, like software developers, conveyed the largest productivity gains from AI. This result is not driven only by coding; it holds when we leave out computer and math occupations. It echoes a previous [Economic Index finding](#) that also favored higher-paid workers: in tasks requiring greater levels of education, Claude tended to reduce the time taken to complete a task (relative to doing it without AI) by a higher percentage.

Some of the lowest-paid workers describe high productivity gains as well. This included a customer service representative using “AI to save me a lot of time with creating a response based on another one.” And in some cases, people in low-wage jobs were using AI on technical side projects. One delivery driver, for example, was using Claude to start an e-commerce business, and a landscaper was building a music application.