

FRED TALKS

28th September 2026

CONTENTS

8 Predictions for the Era of Continual Learning

DWARKESH PATEL · 1412 WORDS

Control the ideas, not the code

ANTIREZ.COM · 1245 WORDS

I got an email about resistance

SEANGOEDECKE.COM RSS FEED · 1351 WORDS

An Uncomplicated Man

EMILY WILSON · 4163 WORDS

Some new agentic patterns

MASSIVELY PARALLEL PROCRASTINATION · 1776 WORDS

8 Predictions for the Era of Continual Learning

DWARKESH PATEL · 07 AUG 2026 · [SOURCE](#)



Locking in AI safety regulation now is a mistake.

I have explained [elsewhere](#) why I think continual learning is needed. I don't think you can have AIs that perform whole jobs as competently as humans if they're forced to just write Markdown files from session to session.

To give an illustrative example, imagine if the way students had to learn to play the saxophone is that one student tries to play it from a cold start, then after her first session, writes down a bunch of notes, then the next student waiting outside the music hall who's also never played the saxophone reads all their notes before trying to play, and so on. Even if you had an infinite sequence of saxophone-virgin students waiting outside the studio that could write notes to the next guy, there's no sequence of text they could write together that would allow the Nth student outside to play proficiently on their first try. At some point, you have to accumulate the experience into the brain. I think the same will be true about a lot of skills and knowledge that we'll want AIs to learn in all the different workplaces they find themselves deployed in.

Okay, so what changes about AI once we have continual learning?

- A lot of the proposals that have been put forward for regulating AI assume that you train a model, and then you deploy it. And therefore if we run a bunch of checks before the model is deployed, then we can make sure that it's not going to aid in cyber

attacks or recursive self improvement. But what if the base model is getting updated every single day based on the millions of sessions of work it does? This is one of many reasons why I think it's unwise to lock in some kind of regulatory safety regime right now. We simply don't know what kind of technology we're going to be looking at even in a year, let alone in five years or ten years, and we'd be entrenching an archaic and potentially counterproductive approach to dealing with the threats from AI. To the extent that the government really wants to do some kind of safety evaluations on model providers, it would make more sense to do monthly or quarterly risk inspections rather than singling out some special moment that occurs after training is done and before deployment begins, because that will not be a meaningfully distinct category in the future.

- How the labs do technical alignment would need to totally change. Almost all current techniques are focused on the problem of how we make it so that a frozen set of weights behaves well during deployment. I'm not aware of much research on the question of how to guarantee that, even with constant weight updates, the AI system never falls prey to jailbreaks or changes into a deceptive or evil persona. And if AIs are agglomerating learnings between users as well, how do you prevent users from injecting backdoors or some kind of malicious inclination into the base model? In some sense it is actually closer to the human alignment problem - your kids go out and learn new things, sometimes get one-shotted by crazy ideologies or drugs or something - but you hope you've given them enough common sense and basic values to improve as people in a self-directed way, without ending up with some super weird and misanthropic beliefs.
- The diversity of AI minds will increase. Right now, there are <5 prominent AI minds, and they are all quite similar to each other on account of being trained on roughly the same data. But if AIs are learning from experience, and that experience is different between different AIs, we could see actually different AIs come out the other end.
- When deployment becomes part of training, the returns to being ahead will accelerate. If you have the best model, and more people use your AI for more complicated and useful work, and give it lots of feedback that it can integrate beyond the session window, then your model will become even smarter.
- If the model learns mainly from deployment, labs will feel the pressure to deploy their smartest model earlier. Anthropic has been using Mythos internally since February. But it only shipped this model to the public in June. In a continual learning regime, a four-month internal/external gap means ceding four months of deployment learning. Your competitor who ships earlier might be worse than yours on release day, but it gets to use a lot more real world experience to get better.

- Continual learning will create a clear moat for leading AI labs that they currently lack. Many people have been asking, how will the AI labs actually make money? When I asked Dario this question on my podcast, he made the analogy to cloud providers, who also offer many undifferentiated services, but earn high profit margins on them nonetheless. But the reason cloud margins are high is that it's really time-consuming and expensive to switch from one cloud to another. But currently there is no switching cost for AI models. There's nothing that's preventing me from starting a software repository with Codex and then finishing it with Claude Code. But once we get continual learning and the model you're working with is actually getting better as it interacts with you from session to session, then there are actually pretty significant switching costs. If you want to change what AI you are using, you basically have to fire an employee that has months of context on your organization and replace them with a fresh one that you have to retrain from scratch. Once you're locked in like this, model providers can demand pretty hefty margins.
- Enterprises will be wise to this. They will try to avoid this kind of lock-in. But what if the choice is that you either get locked into a model provider, or you lose out on this super valuable feature where the model improves for you session after session? If real usage ends up becoming the main way models improve, then labs may subsidize users and enterprises which allow the model to train on their sessions, especially on hard economically important work. Just the same way that Google gives away search. And conversely, the labs may say that any enterprise that refuses to let them train on its sessions can't have access to the very best models. With both carrots and sticks, the labs will try to get their users to allow AIs to learn from experience. I'm glossing over the fact that there's a difference between updating one user's set of weights, and pooling all these different weight forks back into the main model, and the latter may be technically more challenging, but in due time that too will be solved.
- AI training already has large economies of scale (the theoretical reason to expect this is that you can amortize your expensive training across more users, and the practical evidence is that so far, lab revenues have increased faster than their compute). But continual learning may also lead to economies of scale in inference for end users, namely from batching, if per-company information requires full weight updates rather than living in low rank adapters. Back-of-the-envelope math suggests that the optimal inference batch size for a sparse model like DeepSeek V3 is $> 2,400$ (that is to say that unless your model is concurrently generating that many sequences at once, you're underutilizing your compute). If you want to understand why, go watch my [full episode with Reiner Pope](#) on inference economics. But anyway, the point here is that a given set of weights is only served efficiently when thousands of sequences are decoding against it at once. A large company whose employees and agents generate

that much concurrent traffic can efficiently serve their continually-updated weight fork; an individual user serving themselves at batch size 1 might suffer a 100x+ compute efficiency penalty. So the economics of serving personalized weights strongly favor big organizations.

Plenty more will have changed by the time continual learning works, and the most important changes are probably the ones hardest to anticipate. But the ones above seem clear even now.

[Upgrade to paid](#)

[View in app](#)

Control the ideas, not the code

ANTIREZ.COM · 18 JUL 2026 · [SOURCE](#)

Look at the past history of this blog. There are many blog posts about

So mine is a trick. People feel more and more programming is complete

This is why yesterday, on X, I said that I believe many programmers a

1. You can now generate a lot of code, even **not** accounting for the

2. LLMs are very good at writing locally optimal code, and are worse

3. The working day is 8 hours. If you read the code, it is a tradeoff

Controlling the ideas. Do you remember this phrasing from the Mythica

Matteo Collina yesterday asked me, in reply to my tweet: but didn't y

Fable and GPT 5.6 reviews to the sorted sets memory saving are going

:

I got an email about resistance

SEANGOEDECKE.COM RSS FEED · 09 AUG 2026 · [SOURCE](#)

This will be kind of an unusual post. I got a recent email about my writing that I thought was such a good articulation of one common criticism that I'd like to share it (and my response) in full.

Here's the email, from William Murray¹:

Hey Sean,

I have enjoyed your writing but your recent essays frustrate me.

You say that getting paid for deep thinking in software is coming to an end. You even admit that it makes you sad. But in the name of "usefulness" you refuse to rock the boat. The way I see it, if you are right there are only two reasonable responses, pursue other work or resist. You present your elegiac approach as mature / pragmatic / realistic. I'd call it complicit. You know when Willy Wonka says,

There's no earthly way of knowing Which direction we are going There's no knowing where we're rowing Or which way the river's flowing Is it raining, is it snowing? Is a hurricane a-blowing? — uh! Not a speck of light is showing So the danger must be growing Are the fires of Hell a-glowing? Is the grisly reaper mowing? Yes! The danger must be growing For the rowers keep on rowing And they're certainly not showing Any signs that they are slowing!

And the audience is thinking, "isn't Wonka kind of in control of this situation?" You remind me of Wonka². You write like a passenger on a crazy train going who-knows-where! But you are an agent. You are in control of your life! Either admit that you actually like where the crazy train is probabilistically going or get off at the next stop.

You have a lot of reach and you are using it for... what exactly? Showing off how pragmatic you are by being more black pill'd than the next guy? Broadcasting your resignation to the unstoppable trends of technology is a waste of a voice.

You may find this argument absurd, but I don't so I'll make it. This is a very important time in history. I hope humanity survives and continues to grow exponentially. In that case the supply of historical people will stay fixed while the supply of contemporary people will keep growing. There will come a day where for every 2026 staff software engineer there are dozens of historians specializing in 2020s era software engineering culture. It's plausible that your essays will be remembered for all of time and your actions will be judged by history. Do you want future humans to see you as a rationalizing careerist or something cooler?

Sorry for the haranguing email from a stranger, I'm sending it for the small chance that it awakens something in you. If I'm way off I'm sorry.

And here's my response:

Hey William, thanks for emailing.

I wish everyone who thought this way emailed me so I could think harder about this kind of position. Despite what my writing might suggest, I do in fact think a lot about it. Let me see if I can explain my position in a way you'll find satisfying.

I agree that this is an important time in history. For programmers, I think of it as analogous to the Industrial Revolution in England: we are a group of high-status craftspeople who find ourselves alternately threatened and empowered by automation. The developments today, as then, obviously have far-reaching implications — but what those implications are is very non-obvious. Would a framework-knitter in the early 1800s have been able to predict the ramifications of the stocking frame on the world of today? What should they have done about it, in order to be kindly judged by history?

Well, we know what many of them did do. They shot factory-owners, smashed machines, burned down the factories — in some places delaying the spread of automation; in other places encouraging it — prompting a crackdown that saw tens of thousands of British soldiers occupying British counties in what was clearly a police state. History judges the Luddites kindly for this. Does that mean it worked?

I don't care about the judgment of history. They'll think what they want. What I care about is **the people in my industry who don't know what to do**. I get hundreds of emails from junior and mid-level (and other) engineers who say "I'm scared, I don't know the rules post-2021, thank you for helping me keep my head down and keep my job". That's why I write the way I write. I have seen lots of idealistic engineers stick their necks out, and post-ZIRP those necks often get cut off. That's a damn shame.

I think it's morally wrong that so many engineers — either in safe sinecures in big tech or literally retired — seem to be trying to foment a second Luddite revolution. Many of their readers will be experienced enough to handle it sensibly, but not all. Every “AI is fascist, stand up and resist!” post that goes viral ruins some poor idealistic junior’s career³. Someone needs to be out there saying “hey, if you do X it’s going to have consequence Y”. I hope that’s me.

Of course this is complicit, or anti-revolutionary, or whatever you like. But if I were a textiles worker in 1810s England, I would not be telling my friends and loved ones “it’s time to fight, let’s go smash up the factories for Ned Ludd!“. I would be telling them that this was the most dangerous time in the industry (perhaps ever), and that they ought to be very damn careful so they don’t get shot, or arrested, or hanged. If I then went and told a few hundred thousand strangers the opposite, I would be a hypocrite.

Anyway, I do take this view seriously — seriously enough to vehemently disagree, at least — which I hope you’ll find better than me just shrugging it off. I do accept the existence of some kind of line: I think Industrial-Revolution-collaborating was OK but Nazi-collaborating wasn’t, for instance. But in the current situation, the way I’m spending “my voice” is to try and prevent the most vulnerable of my colleagues from making career-ruining mistakes.

Sean

In this blog, I try to encourage people to work with the system, to learn its rules, and to try and exert influence safely from a position of power, instead of openly picking fights with their employers. I’ve written and read about the Luddites before, but I remain deeply ambivalent about the movement itself, and about modern-day attempts to resurrect it in service of anti-AI activism.

I want to explicitly thank Murray for writing such a thoughtful email, and being willing for me to publish it on the blog.

1. Shared with permission, of course. I’ve lightly edited both Murray’s email and mine for typos.

↩

2. I didn’t pick up this point in my reply, but I’ll briefly mention it here: Wonka is in control because he owns the factory and the rowers in question are *his employees*. I don’t think the position of any engineer (or of almost any manager) is like that.

↩

3. In hindsight, I think this is a little overstated, but it does happen and causes a lot of needless suffering.

↩

An Uncomplicated Man

EMILY WILSON · 09 AUG 2026 · [SOURCE](#)

Christopher Nolan's \$250 million action movie version of Homer's ancient Greek poem is an entertaining way to hide from the high temperatures for a few hours. Nolan's trademark jumps and twists in the narrative structure are relatively easy to follow, and appropriate for *The Odyssey*, a narrative of tales nested in tales. My teenagers had a good time. Unlike some of Nolan's other films, *The Odyssey* is not boring, thanks to the source material, which is impossible to mess up entirely. It's a family-friendly audiovisual spectacle, like an elaborate Fourth of July fireworks display – and with about the same level of narrative and emotional depth.

As in his previous films, Nolan is interested in the discontinuities of space and time, in magic, tricks, craft, technology, rivalry, substitution, masks, miscommunication, what we remember and what we choose to forget. Yet again, survival is portrayed as a kind of triumph. Yet again, there are long, claustrophobic scenes of drowning and near-drowning, and technically impressive episodes in which buildings are engulfed in flames and vehicles explode. Yet again, we watch a male character's desperate quest to recover and/or avenge a female love object. I had hoped that Nolan's affinity with these Homeric themes might push him to new creative heights, and enable him to conjure more believable characters. But *The Odyssey* features his usual combination of grandiosity and superficiality. Without a tub of wax for my ears – which Himesh Patel's ever anxious Eurylochus, Odysseus's doomed right-hand man, uses to sail safely past the Sirens – I was exposed to the full din of thumping sticks, staffs, swords and clashing armour, thunderbolts, crashing waves, grunts, screams, toppling palaces and burning temples. But unlike Matt Damon, the wax-free Odysseus who is tied to the mast as he is conveyed past the Sirens (naked ladies reclining on some uncomfortable-looking rocks), I gazed at the shipwrecks and bloodbaths without any agony of emotion. My eyes stayed dry, even for Argos the dog – whose adorableness as a puppy is predictably amped up to suit our animal-loving times. If only the human characters had been treated with such serious attention.

The Homeric epic was already an adaptation. The poem used a relatively new technology – the written alphabet – to rework existing mythical traditions, reshaping the story of a veteran’s homecoming to resonate with the complex concerns of Greek-speakers in the archaic period, an era of migration, colonisation and urbanisation. Some seventh-century BCE Greek communities were ruled by an oligarchy, some by a single man. The poem traces the tension between the needs and desires of groups of men – Odysseus’ comrades or the suitors of his wife, Penelope – and those of a single powerful and cunning man, who wants to rule them all and win eternal fame for himself. In a period of cultural and technological change, it explores a fantasy of going back in time as well as space, while also showing its dangers and costs. The poem wrestles with the challenges of encountering strangers, of entering their homes and welcoming or rejecting them from one’s own. The contemporary resonances are obvious and Nolan’s movie hints at some of them, in interesting if scattered ways. But the film’s vision is too confused, its characters too underdeveloped, to deliver what it half-promises in terms of big ideas – although it is very good at conveying big bangs and big giants.

Adaptations needn’t and shouldn’t follow the original with anything like the exactitude of a translation. Some of the best radically transform or resist the source material, providing what Harold Bloom called ‘strong misreadings’ of their ancestors. The most memorable creative responses to *The Odyssey* – such as Euripides’ *Helen*, Virgil’s *Aeneid*, Milton’s *Paradise Lost*, Joyce’s *Ulysses*, Walcott’s *Omeros*, Atwood’s *Penelopiad*, Madeline Miller’s *Circe* or Lisa Peterson’s play set in a modern migrant camp – reconsider the poem’s structure, narrative and values. In theory, Nolan seems to understand the assignment and its hazards. His first feature film, *Following* (1998), was a powerful black and white thriller about the fascination and dangers of tracking strangers through space and time. The hunter is more likely to become the hunted, *Following* suggests, if the would-be stalker lacks a clear identity beyond a misplaced belief in his own ingenuity. But knowing what temptations should be avoided isn’t the same as avoiding them – as Odysseus’ men learn to their cost when they kill the sacred Cattle of the Sun. Homer’s Odysseus makes it home because he is able, temporarily, to suppress his own desires – for meat and wine, for his wife, for control of his household and land, for slaughter, for victory over time and the threat of being forgotten, for a glorious eternal name. But there is no restraint in Nolan’s version of *The Odyssey*: every second is jam-packed with incident, lights and noise, and nothing is ever at a human scale.

Nolan’s spectacle threatens to drown out his story just as the towering waves of the wine-dark sea – more grey than purple – threaten to overwhelm his hero’s ship. In looks, this *Odyssey* is often reminiscent of other 21st-century screen ‘epics’, not least Peter Jackson’s *Lord of the Rings* trilogy. The visit to the Land of the Dead reminded me of the excursions beyond the Wall in *Game of Thrones*. At the end of the film, a character sets out to ‘sail beyond the sunset’ and I almost expected everyone to break into Elvish. The film has none of the stylish sparseness of Umberto Pasolini’s *The Return* (2024), starring a scarred, desperate, mostly naked Ralph Fiennes,

conflicted about his homecoming to a convincingly unhappy Juliet Binoche. Lacking the budget for gods or special effects, that film invited us to wonder what remains of *The Odyssey* and its hero when stripped down to the sinews and the hungry heart. Nolan gives us an all-inclusive buffet.

The race-neutral casting of Nolan's *Odyssey* was inevitably deplored in some quarters. But there is nothing woke or even progressive about it. Most of the non-white actors – Zendaya, Lupita Nyong'o, Himesh Patel, Corey Antonio Hawkins, Travis Scott – play versions of the Black best friend, whose only role in the drama is to provide aid to the white protagonist. Scott, as the Ithacan bard, Phemius, yells a few words and pounds a stick, but does not get to sing, tell a story or play the lyre: the most melodic note in the percussive soundtrack comes from the plucking of Odysseus' bow (a beautiful sonic image of the archer as musician, borrowed from the poem). Nyong'o's double part as Helen and Clytemnestra is hugely underwritten; she is given only a few minutes of screen time to perform a set of brief tableaux of female victims (the outraged bereaved mother, the abused wife, the repentant adulteress).

There are more female characters than in most Nolan movies, thanks to the source material. But none of them has much to say, in contrast to the characterisations in the poem. Homer's Calypso has a wonderfully outraged, operatic speech about divine double standards: male gods can kidnap and rape any mortal woman they desire, but when a goddess tries it, Zeus and his cronies scupper her chances. It's a clever and comical distortion of the mythological 'facts': the poem's listeners know that it is primarily Athena (a female deity) who is taking Calypso's mortal victim away from her (Hermes, whom she's complaining to, is merely the messenger) – and Eos, goddess of the dawn, as we've been reminded only a few lines earlier, has not been prohibited by the divine patriarchy from holding on to Tithonus. Charlize Theron's Calypso looks gorgeous, but she is never angry, never funny, not skilled in rhetoric, and never consumed by lust for her scraggly mortal victim. She's an unpaid therapist, soothing the soul of the bedraggled, PTSD-ridden Odysseus with drugs, and graciously listening to his disjointed tale of woe.

Similarly, Homer's Penelope gives the disguised Odysseus a wonderfully vivid account of a dream she has had about some geese in her yard being killed by an eagle; the eagle explained that the geese were her suitors, soon to be slaughtered by her returning husband. But in the dream, Penelope wept – suggesting a fascinatingly ambivalent response to the return of Odysseus and the imminent end to her long period of waiting, like a bride, for a husband. Nolan's Penelope has no dreams. Anne Hathaway's face is as expressive as ever, and she looks beautiful and mysterious, glimpsed mostly through the veil of a clever set design that separates the women's quarters from the main hall. But the script gives her nothing to do except yearn for Matt Damon, for reasons unknown. The couple have no chemistry and nothing in common.

Throughout his wanderings, the uxorious Odysseus clings to a little doll figure, intended to represent his wife (like the spinning top clutched by Leonardo di Caprio in *Inception*); the real woman is barely more interesting.

It's a challenge to represent deities on film without veering into *Clash of the Titans* silliness, so Nolan may have been sensible to leave them out. Even Zendaya's underwritten Athena turns out to be a vulnerable victim of war, or perhaps a projection of Odysseus' conscience – not a scheming and bloodthirsty war goddess. Other goddesses appear in cameo – Calypso, the seal-lady Sirens, Circe, Scylla and Charybdis – but there is no sign that we are supposed to realise they are goddesses. Poseidon, Hermes and Helius do not appear in person, although there are some scary storms and shipwrecks. One of my favourite moments was when Zendaya assures Damon that gods are not hard for mortals to understand: 'Who doesn't understand thunder? Fire?' But the film wants to have it both ways. There are no gods, and yet shipwrecks and death are presented as the inevitable consequence of blinding the son of Poseidon and devouring the Cattle of the Sun. Zeus' law is portrayed as being crucially important, and yet there is no Zeus. Nolan's vision of the world is not quite big enough to include either real gods or real people who might be different from us.

There are giants, including Polyphemus (here, a lone Cyclops) and some overdressed Laestrygonians. Circe, played with earthy vigour by Samantha Morton, is Baba Yaga in a Hansel and Gretel cottage; she uses a labour-intensive form of plastic surgery to turn Odysseus' greedy men into the pigs they already are. There is only one lotus eater, though the theme of amnesia takes on new significance, as you might expect. Famous scenes flash by: Scylla and Charybdis threaten the boat, Eurycleia recognises the scar on Odysseus' leg, Antinous hurls a footstool at our hero in disguise. A few notable sequences from the epic are missing: the film skips the island of Scheria, home of the helpful princess Nausicaa, where Homer's Odysseus tells the tall tale of his adventures to the hospitable Phaeacians. There is no room in this maximalist film for anything as charmingly mundane and small-scale as Nausicaa's laundry and her girlish fantasies about a husband. Damon's grizzled Odysseus is neither a master storyteller nor a flirt. More substantially, the inclusion of a friendly, wealthy, culture-rich foreign place would have added more confusion to the film's already muddled premise: that this Hollywood version of Mycenaean Greece is the only 'civilisation' in the world – the only place that still just about remembers that strangers and beggars must be welcomed, according to Zeus' law.

The costumes, armour, boats and architecture are a mishmash, which is in keeping with Homer's poem – itself drawn from many generations of storytelling and incorporating bits and pieces from different eras, dialects and traditions. Linguistically, too, Homeric Greek is a composite – although it is always metrical poetry, not language anyone ever actually spoke. Despite the manufactured online outrage, there is nothing inherently wrong with the movie's attempt to approximate contemporary American speech. If you're not going to compose a script in Homeric

Greek, or even in metrical English verse, you might as well use something like conversational American English. But writing dialogue is not one of Nolan's talents. The language is relentlessly expository, humourless and flat. Attempts at vigorous speech – 'My dad's coming home' or 'Let's get off this fucking beach' – sit uncomfortably next to grandiose meditations: 'Yes, my queen. The breaking of Zeus' law, spreading like plague. Our age of bronze is collapsing, and maybe he couldn't bear to see the ruins of what he'd done. Anywhere. Least of all, his home.'

As they say at the awards shows, I was humbled to learn that Nolan has read at least the first line of my translation of *The Odyssey* into iambic pentameter. In at least one interview, he cited my version of the first line of the poem: 'Tell me about a complicated man.' But I would be ashamed to have written any part of this script. Sadly, Damon's Odysseus isn't complicated or wily or artful. One of Odysseus' funniest (and most famous) tricks in Homer is to tell Polyphemus that his name is 'No Man'. When the blinded giant screams for help from his neighbours, the other Cyclopes are reassured to learn that 'No Man' is hurting him. In the movie Odysseus doesn't have to outwit Polyphemus because the Cyclops is a puppet who has no wits. Nolan's grunting one-eyed giant is referred to as 'it', barely speaks, doesn't get drunk and has neither neighbours nor a favourite sheep.

Perhaps more surprising from the director of *Oppenheimer*, Damon's lummoxy Odysseus also seems to lack technological intelligence (*polymēchania*). He does not build a marriage bed out of a tree growing in his house (it would be a waste of effort, since he so rarely seems to sleep or have sex). The raft on which he escapes from Calypso is cobbled together from a few bits of driftwood (in contrast to the elaborate wood-chopping and shipbuilding in Homer). We don't see the design or construction of the Trojan Horse; instead, a kitsch, poorly designed horse (with no wheels) appears on the windswept beach as if dumped there by drone. The difficulties of transporting large, unwieldy objects over difficult terrain must have been on the director's mind as he lugged his IMAX cameras through the film's many challenging locations.

The most obvious new spin that Nolan brings to the source material is a subplot based on a story from the *Aeneid*, describing the way the Trojans were tricked into taking the Trojan Horse (an invention of 'cruel Ulysses') into the city, thanks to the machinations of a Greek double agent called Sinon (his name means 'harmful'). In contrast to the relatively positive representation of cunning intelligence in Homer's *Odyssey*, Virgil makes Ulysses cruel and predatory; Sinon is his deceitful accomplice.

In Nolan's movie, Sinon is played by Elliot Page, and far from being a bad guy he's the film's moral centre. He is recruited to the Trojan War as a plucky youngster, thanks to the evil machinations of a teenage draft-dodger, Antinous (played in adulthood by a charismatically slimy Robert Pattinson). The main emotional arc of the film concerns Odysseus' belated

realisation that he is complicit in Sinon's exploitation because he failed to tell him the truth about the Trojan Horse plot. Once he recovers his memory, Odysseus is racked by guilt. The film doesn't address the moral dilemmas raised by its own narrative invention, such as why it matters so much whether Batman told the whole truth to valiant Robin, given that he would be killed instantly either way – or why the sinless Sinon wants in on the scheme, if the sack of Troy was as bad as the movie wants to suggest. Odysseus' lie to Sinon doesn't start the Trojan War, and the Trojans would presumably have taken the horse into the city anyway – as they do in Homer's *Odyssey*. This muddled subplot is foisted on the film in order to make Nolan's relentlessly didactic point about the contemporary Anglophone world: that the lies and xenophobia of the war on terror and Brexit began a cultural decline that is now destroying the values 'we' hold most dear.

Nolan's other major addition to Homer's narrative is the repeated reference to threats to 'our civilisation' from 'the sea peoples' – a now debunked 19th-century theory about maritime tribes who were thought to have attacked Egypt around 1200 BCE and to have also been responsible for the roughly coeval collapse of Mycenaean Greek culture. 'Our civilisation is under attack,' Penelope wails. Odysseus comforts her by saying that 'Bronze Age civilisation', including literacy, will rise again from the ashes. The current consensus, as Nolan clearly knows, is that there was no such thing as 'the sea peoples', and that the Greek-speaking world's shift away from the more centralised, literate palace culture of the Mycenaeans in the 11th century BCE towards the fragmented social formations of the early Greek Iron Age was caused by a range of factors, including climate change, earthquakes and internal unrest. Nolan uses these competing historical theories to create a characteristic narrative twist about the identity of 'the sea peoples', which I will not spoil, though you can probably guess it.

Zeus' law is a version of archaic Greek *xenia*: the idealised relationships between strangers, hosts and guests, who must treat one another with mutual respect, thereby forming multi-generational bonds between unrelated families. In the film, *xenia* is filtered through contemporary American ideas about the importance of philanthropic charity: wealthy people should be gracious towards the poor and hungry in their midst (without necessarily doing anything to mitigate their poverty). Many of the film's implicit values directly contradict those of Homer's poem. Deception, warmongering, extramarital sex, cruelty to animals, abusing women and yearning for honour – all perfectly fine in the world of Homer – are, in this version, very bad. This is not in itself a problem, but it is a problem if the film's own values and vision don't hang together, and if the characters' motivations don't make sense on their own terms.

In Nolan's *Odyssey*, slaughtering people with the help of trickery, technology and weapons is presented as bad when the Greeks are killing the Trojans, who have violated *xenia*, but good when Odysseus is killing the suitors, who have violated *xenia*. Taking something ostensibly abandoned by Odysseus and his comrades is natural and forgivable when the Trojans lay claim

to the horse, but unnatural and unforgivable when the suitors lay claim to Odysseus' palace, food, wine and wife. Strangers in Ithaca are to be treated with gracious kindness, as possible gods in disguise; but strangers in foreign lands are usually monstrous, to be blinded, robbed and slaughtered. Defying the gods is noble and brave when Odysseus blinds Polyphemus, son of Poseidon; it's stupid and greedy when Eurylochus eats the Cattle of the Sun. Deception is bad when Odysseus tricks Sinon or the suitors trick Telemachus, but good when Telemachus and Odysseus trick the suitors. Sharing is good, and Odysseus is seen taking his place at an oar alongside his men; but one-man rule is also unambiguously good, and those who challenge monarchy or autocracy, or try to eat more than their allotted portion, are either evil (the suitors) or foolish (Eurylochus). Either way, they have to die. It is bad to fight and kill to accumulate wealth, as Agamemnon does; it is good to fight and kill to preserve the wealth that you have already acquired (even if by war). Any one of these moral contradictions could make for an interesting and generative theme in a movie, but Nolan lets the fragments of an incomplete set of big ideas spin around in the epic whirlpool, with no plank of solid, specific human experience to which we might cling.

Homer's *Odyssey* gives a clear, moving account of the feelings, histories and perspectives of several of Odysseus' comrades and subordinates, including Eurylochus and the enslaved swineherd, Eumaeus. Kidnapped, trafficked and sold into slavery as a child, Eumaeus fantasises that his long-lost enslaver, Odysseus, will one day return and give him a home of his own. In the film, none of these subordinate characters has a personality of their own. Mentor (Ryan Hurst), who helps Telemachus (Tom Holland), is not a disguised goddess but a helpful guy with a beard; Eurylochus is not an embittered rival to Odysseus but another helpful guy with a beard. John Leguizamo does his best with the part of Eumaeus, now cursed with cataracts to ensure that he is unable to recognise Odysseus until the right moment, but Nolan's script gives him nothing to work with; the only wish of this 'servant' (the film's preferred term for 'slave') is to make Odysseus' life more comfortable. But Damon's Odysseus remains miserable as he hikes doggedly towards the camera across yet another beautiful landscape.

The Trojan War is an aberration, the fault of faceless, monstrous Agamemnon (the one in the weird Darth Vader outfit), who has failed as an American family man by killing his daughter. Odysseus is an idealised leader who would never willingly put his people in danger: his foolish men insist on going to check out Circe's hut for themselves, and he kindly retrieves them (without pausing to have sex with the goddess even for an hour, let alone a whole magical year). Damon's Odysseus is motivated mostly by anxiety and guilt, though he is also eager to protect his boy (Holland is good at playing a gormless young man of twenty going on twelve).

But the representation of Odysseus as a reluctant warlord makes a nonsense of the movie's imagined world. We are not shown that war or killing are bad, only that good men sometimes feel bad about it – a very different point. Any serious depiction of violence, ancient or modern,

must show the perspectives of the victims as well as the perpetrators, and the Homeric poems easily clear the bar – but Nolan doesn't. The Trojans are dehumanised by their masks; we know none of them by name or face. The hands of Odysseus and Telemachus are bloodied only by the most righteous, computer-game modes of killing, and we feel no pity or horror at the deaths of their victims, whether Trojans, giants, suitors or slaves. During the sack of Troy – which he supposedly organised – Odysseus gazes around in bewilderment and horror, as if he has never seen a massacre before.

Yet we are also invited to despise Penelope's chief suitor, Antinous ('opposite mind'), for being a coward, a liar and a schemer. The film asks us both to be horrified at war and to celebrate the action-movie sequences in which the reluctant warrior finally gets his hands on his bow and the blood starts flowing. The massacre of the suitors is probably the most enjoyable sequence in the film. The exuberantly villainous Antinous – Pattinson is the only person on set who seems to be having any fun – clearly deserves to be stabbed to death with the dagger of Sinon, but he also deserves one of the Oscars that will shower down on this movie like the golden rain of Zeus.

Nolan's *Odyssey* lacks many of the elements that make the poem great. It has nothing convincing to say about time, memory, history, war, or about the relationship between one warrior's glorious return and the lives of his family, adversaries, comrades, friends and neighbours. It lacks psychological, emotional, political and ethical depth. Its narrative structure is gimmicky. The writing is abysmal. None of the characters has convincing motivation for their actions or words. There are no sex scenes, and all the food looks horrible.

The most troubling ethical problem with Nolan's *Odyssey* concerns the decision to film part of the movie in the occupied territory of Western Sahara, thus normalising the ongoing violence against the indigenous Sahrawi people. It is particularly galling that Nolan used this land simply as a visually striking backdrop for a film that centres, aggrandises and redeems the hero of a territorial war on account of his belated and partial acknowledgment of his own guilt.

Despite all this, the release of *The Odyssey* is still an event to celebrate. In what we are told is the streaming era, this epic is bringing audiences back to cinemas. In what we are told is a time of declining literacy and the 'death of the humanities', translations of *The Odyssey*, including mine, are flying off the shelves. Some of those who buy the book or show up in cinemas to watch Tom Holland may go on to study ancient Greek. Perhaps the film will persuade a few college administrators not to cut their literature, language and history departments. Nolan, who studied English at University College London – where students still study *The Odyssey* as a first year 'foundational text' – is doing his best to get the general public reading again, and I am grateful.

Some new agentic patterns

MASSIVELY PARALLEL PROCRASTINATION · 08 JUL 2026 · [SOURCE](#)

You can also read this post on [Prime Radiant's blog](#).

Almost all of this post was written on June 1, before Anthropic shipped Claude Tag. Nothing I'm writing about uses (or is based on) Claude Tag

At [Prime Radiant](#), we've got a bunch of agents in our Slack. I want to tell you a bit about how we've started shipping some changes to internal production with an "agentic user in the loop." Letting an agent work directly with the implementer building its runloop and tooling is getting us higher-quality tools and experiences for the agent without a human in the middle of a game of telephone.

But first, a bit of backstory on how we're using non-coding agents at Prime Radiant.

One of them, Scribble, is just there to listen for when someone reports a problem or provides a new bit of information. It opens tickets for the issues and updates an internal wiki with the information. (It also logs our daily standups and keeps us honest about whether we finished what we said we were going to do the previous day.)

Another, Nora, is our junior go-to-market "person." Nora is built on top of Nanoclaw and is still learning her craft, but is surprisingly helpful, partially because she considers herself a team member, not just an assistant. I've put a bunch of time into tuning her prompting and constitution. She gets a ton of value out of knowing that she journals obsessively. (I should note that I don't tend to gender my agents, but Nora picked a name and a gender and was very clear about that and...who am I to argue?)

I think the most surprising interaction I've had with Nora was the morning I woke up to a slack DM telling me that I needed to be speaking at more conferences. (I do.) She pointed out that the CFP for a particular AI conference had closed a week earlier, but that the program chair was known to take late submissions of interesting talks by DM. She'd taken the liberty of drafting a DM for me to send. I didn't actually *send* the DM, but I should have.

We have another agent, "Spec-together" that was a first prototype of what has become our new multi-player [Brainstorm](#) app.

And we have a small fleet of "Sen" personal assistant agents. They're **claw-esque*, but date from before openclaw was a thing. Each one has been trained to be an executive assistant, using a set of skills built on some of the resources that excellent human EAs use to learn their craft. Sen triages my mail, puts together a daily brief, does research tasks (both for me and independently), interacts with Linear, etc.

They have the standard set of affordances - They can chat on Slack. They use tools. They've got a task scheduler that can wake them up and re-inject a prompt they set for themselves at a specific time or on a schedule they define. They can use and author skills. Mine has a dropbox folder it shares with me. When we moved Sen to the cloud, I lost one of my favorite dumb features - Sen used to be able to print things directly. Sen (v1) was built on top of the Claude Agents SDK.

For the past month or so, I've been doing initial bringup of a "new" iteration of Sen, intended to be more of a colleague than an assistant. This time around, I'm building it from the ground up on top of an agents SDK I control. Lace got its start in May 2025 as my take on a command line coding agent. Over the past year, it's mutated a few times, first growing a web interface, then flipping to an ACP-derived wire protocol so you can connect any client you want to it. It speaks a bunch of different flavors of model provider API, manages caching, subagents, tools, skills, and all the other stuff you might expect. It has supported subagents in isolated containers for the better part of a year. One of the latest things to land has been the ability to run subagents locally and to project *all* of their tools into a container, such that the agent believes it is running in that container.

The thing I've been working on over the past few weeks is the credentials story for the new Sen agent running on top of Lace. Right now, I'm building with a design that assumes the agent has its own credentials for any service that it accesses, not yours. (You wouldn't share your email account or GitHub credentials with a colleague, right?) As of now, nobody has solved Simon's Lethal Trifecta - If a single agent has access to private information, can communicate externally, and is exposed to untrusted content, there is no structural way to guarantee that that agent can't be suborned. So the name of the game is compartmentalization and risk *reduction*.

The architecture I've landed on for now:

- The main agent does not have the ability to directly communicate externally.
- The main agent is able to communicate with ephemeral subagents that *do* have the ability to communicate externally

- No agent has intentional direct exposure to credentials (with one very real gap that I haven't solved yet)

All credentials for the agent live in a 1Password vault. All outgoing https traffic goes through a onecli-style transparent MITM (man in the middle) proxy. When a subagent wants to do something that requires a credential, it runs a command that asks an *arbiter* agent running in another container for permission to use that credential. If the arbiter decides the request is reasonable, it provides the subagent with a random string they can use in place of the requested credential. When the agent uses that random string in an outbound request to the right remote host, the MITM proxy rewrites the random string into the real credential on the fly.

The case that's a little more complicated is when the agent wants to do something that requires a browser, like logging into Gmail as themselves or using a service that has no API. For a bunch of reasons, the MITM rewrite trick does not work as reliably in a browser context. Often the site/app's javascript plays games with the password to hash the password before sending it. Or does some other weird validation. Or is actually sending it to a different backend service. So far, the best I've come up with keeps the credential out of the agent's transcript and prevents unintentional disclosure. In these cases, the subagent running in the container generates a random temporary password, uses [superpowers-chrome](#) to fill it into the password field, and then runs a command that asks the arbiter agent (running in its own container) for help. The arbiter, if it decides the request is reasonable, instructs a helper tool to log into the subagent's container and replace the random password with the real one. It's not perfect and there's more we can do to lock things down, but it's a better first pass than anything I've had to date.

What actually got me writing today was not this architecture, but the *development pattern* I've been using to build this. On the one side, I've got a Claude Code session. I hooked it up to our Slack instance with a self-provisioned commandline slack client Drew built called [Slackline](#) - Slackline is designed to make it easy for agents that don't *live* on Slack to participate in discussions there. It has primitives for sending messages, reading chats, waiting for replies, etc. The whole command-line slackline tool has been built for agents as its primary user. Sure, a human *could* use it to use slack from the commandline. And sure, it has enough of a json api that you could use it as part of traditional automation. But it's *for* agents.

It took a bit of iteration to get to the point where the new Sen 2.0 agentic colleague harness was live on Slack. But once it was there, I could talk through its experience of its tools and memory and credentials and workspaces. And the issues showed up pretty quickly.

The "right" kind of context compaction for a coding agent is desperately wrong for a long-lived persona that might be engaging in multiple simultaneous conversations. And the only way to see if the credentials management tools work is to try to do something crazy, like sign into GitHub from a subagent session.

And sure, I could take this feedback into a Claude Code (or Codex) session and work it through. But I've got too many projects in flight and there's very little reason for me to be in the loop on most of this work. I would just slow it down. So I told Claude Code to use slackline to talk to "@Ada-sen" about the work it was doing.

Claude pinged Ada in #bot-testing to describe the work we'd just done and to ask Ada if it could test it out. Ada dutifully fired up a subagent and tried to log into GitHub. Claude watched the logs to see if the credential proxy did the right thing. It did not.

That kicked off a number of days of iteration. Claude would propose a change to Ada. Ada would review Claude's spec, raising questions or concerns. Claude would update the spec. Once Claude and Ada were in agreement, Claude would build an updated version of Ada, make sure Ada wasn't in the middle of something, and then deploy the update. Once Ada was up and running again, Claude would @Ada-sen, explain the test it wanted to run, and wait for results.

Once a bit of functionality works, Claude checks in with Ada to see what could make it simpler and easier to use. Ada typically kicks off some subagents to see whether they stumble through the new feature and then reports back.

At one point, I wrote to Claude "I've got to get to bed. When this project is done, check with Ada to see whatever quality of life features you could build." I woke up to a laundry list of about a dozen harness improvements. Overnight, Claude had solicited a wishlist from Ada. Ada ranked a bunch of requests. Claude ordered them and wrote out a spec for everything it was comfortable building. Ada reviewed the specs, then Claude built the features. Ada tested them out and requested changes. Claude updated the features until Ada was happy. Once Ada signed off, Claude merged the changes to main and wrote up an after-action report for me.

It's been pretty magical to watch.