FRED TALKS

22nd October 2026

CONTENTS

The Self-Help Trap: What 20+ Years of “Optimizing” Has Taught Me

TIM FERRISS · 3114 WORDS

An Alien Mind | OpenAI

OPENAI · 3161 WORDS

The Self-Help Trap: What 20+ Years of "Optimizing" Has Taught Me

TIM FERRISS · 08 MAR 2026 · SOURCE

One danger of modern self-help.

> **"We cannot reason ourselves out of our basic irrationality. All we can do is to learn the art of being irrational in a reasonable way."**
> — Aldous Huxley, *Island*

It was cold out, but none of us were cold.

I sat with five men in the mountains of Montana. As the sun set, the fire in the center cast dancing light on our faces. Reclined against fallen trees in a tight circle, we ate mushrooms and fish we'd found under trees and along streams. The whole crew burst into laughter yet again, and one of the guides passed around a fresh batch of pine needle tea.

Bathed in warmth, I took off a layer and glanced skyward through an opening in the trees. The stars shone like crystals on black velvet, and the show—the biggest meteor shower of the year—was starting.

In that moment, there was nothing to do. Nothing to improve. Nothing to fix.

It was perfect.

***

The older I get, the more I think that self-help can be a trap. Sometimes the cure is worse than the disease. I say this after ~20 years of writing self-help and a lifetime of consuming it.

Spend enough time in the world of "improvement," and you'll notice something strange: The people most obsessed with self-help are often the least helped by it. Behind the smiles and motivational quotes, behind closed doors and after a drink or two, the truth is that they're not able to outsmart their worries.

On one hand, perhaps this unhappiness is precisely what lands one in self-development in the first place, right? I long assumed this about myself, and it's partially true.

On the other hand, what if self-help *itself* is actually creating or amplifying unhappiness?

Modern self-help contains an in-built flaw:

To continually improve yourself, you must continually locate the ways you are broken.

Fortunately, there are a few perspective shifts that make all the difference. It took me embarrassingly long to figure them out.

To get started, let's take a fresh look at an old concept.

MASLOW'S HIERARCHY OF NEEDS?

> **"I suppose it is tempting, if the only tool you have is a hammer, to treat everything as if it were a nail."**
> — Abraham Maslow

Maslow's Hierarchy of Needs has captured the minds of hundreds of millions. It offers simplicity in a terrifyingly complex world.

Abraham Maslow's "A Theory of Human Motivation" (1943) contains five levels, which are typically presented like the below pyramid. This one is pulled from the Wikipedia entry on the subject:

We've all seen it. Clear as day, you can see the goal post at the top: self-actualization.

LFG! It's time to journal and 80/20 myself! Pass me a shaman and some modafinil.

That's the mission. That's the point.

Right?

But hold on. A critical footnote got lost in the shuffle. In his later writings, especially notes compiled in *The Farther Reaches of Human Nature* (1971), Maslow added a *sixth* level above self-actualization:

Self-transcendence.

That update never quite made it out of the crib. The consultants are to blame, but that comes later.

Self-transcendence means going beyond the self — seeking connection with something greater, such as service to others, nature, art, or the divine. Why is it important? Well, for one thing, as Tony Robbins put it at an event long ago: "'I, I, I, me, me, me' gets to be a really fucking boring song."

But it's not just a boring song; it's dangerous to your health.

DON'T BE A SOMO

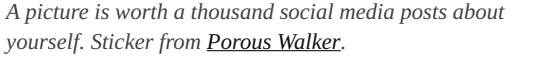
> **"The man who renounces himself, comes to himself."**
> — Ralph Waldo Emerson

Self-help is dangerous precisely because it easily becomes self-fixation.

A focus on improving the self usually first requires finding problems with the self. This is quite the pickle. In a society that rewards problem-solving, you can end up hallucinating or exaggerating unease in order to fix it. This leaves you always in the red, always one step behind. Imagine a dog chasing its tail that has committed to being unhappy until it catches the tail… but it's always just a few inches short. Still, it whirls around and around, "doing the work." Perfection always recedes by one more book, one more seminar, one more habit tracker.

Put in more colorful terms, misdirected self-help turns you into a self-obsessed masturbatory ouroboros (SOMO).

To remind me of the SOMO risk, I have this sticker on my laptop:



A picture is worth a thousand social media posts about yourself. Sticker from Porous Walker.

Now, to be clear, I still love self-help. Ain't no way Timmy can give up the sauce. There's a place for it.

From The Bible to Seneca, and from Ben Franklin to Stephen Covey and far beyond, there's a lot of valuable advice worth taking. I used to mainline it all – no time to waste! – and jump straight into action. This did some good, but there was a lot of collateral damage.

Why?

Because there are at least three "tectonic plates of self-help" that I couldn't see for decades, and they dictate how much net-positive or net-negative comes from all the striving. Before you sprint, you want to calibrate your direction.

THE THREE TECTONIC PLATES OF SELF-HELP

> **“As to methods there may be a million and then some, but principles are few. The man who grasps principles can successfully select his own methods. The man who tries methods, ignoring principles, is sure to have trouble.”**
> — Harrington Emerson

In the last few years, my life has become much more of a joy than a grind, and that’s because I’ve focused on three tectonic plates.

Let’s take a close look at each.

1. Intention

Individual or Social?

Americans, in particular, worship at the altar of the rugged individualist. There are clear upsides to this. But steeped in a culture – offline and especially online – that puts the self on a pedestal, we can take self-improvement to be an end unto itself: a better self.

But is it an end unto itself? Does it automatically produce good things? I now have my doubts.

Here’s one analogy I’ve drawn for myself.

Let’s pretend that life is the game of soccer. You can work on the mechanics of soccer by yourself. You can always get better at dribbling, shooting, and running drills as a solo practitioner. You can read dozens of books, study tape, and earn a PhD in the physics of ball flight. You can post videos of stunning shots on YouTube and get showered by emojis.

But none of this is actually playing the game of soccer.

You can spend your whole life preparing for, instead of playing, the game of life.

But why would anyone, including yours truly, succumb to this?

Subconsciously, it spares you from the messiest but most rewarding game of all: human interaction. Perhaps people hurt or traumatized you long ago. You might also justify the endless polishing, as I did, with some version of “Once I’ve perfected myself, then I’ll be ready for relationships.” But here’s the rub: that practice is exactly endless. You can always get better at dribbling and penalty kicks.

Digging further, focusing on improving the self is often in service of trying to control the world, especially if things were unpredictable or unstable when growing up. Banish emotion, live by spreadsheets, and all can be well. All can be *controlled*, or so the illusion goes. But as soon as you're interacting with – let alone depending on – other people, control as a construct goes out the window. And so we consciously or subconsciously avoid the messiness. This is also one of the reasons why a lot of optimizing achiever folks have a hard time in intimate relationships.

So how do I think about "self-help" now, having realized all of the above?

It is refreshingly simple: the goal is to build and improve my relationships. The sooner you get on the real field with real players, the sooner you can get to playing soccer and engaging with life. No more auto-fellating, even with the best of intentions. We've evolved over millions of years to be deeply social creatures, and the more you dodge that IN REAL PHYSICAL LIFE, the more you will suffer. This is why solitary confinement in prisons is often considered cruel and unusual punishment… and yet we do it to ourselves all the time.

There are a few questions that help corral this tectonic plate of *intention*:

- How does any given "self-help" help me in my relationships, and how can I apply it with other people *today* or *this week*?

- How can I take the ship out of the harbor and test it where it counts?

2. Audience

Do you have an audience for your self-development? If so, be careful.

Nary a minute can be spent on social media without bumping into a CAPS-rich "HOW X CHANGED MY LIFE" or a photo carousel of an ayahuasca retreat. If only Costa Rica got a dime for every bikini-clad healer under a waterfall!

Welcome to the theater of performative self-help. I won't belabor this, as we've all seen it, but I suggest reading about the insidious creep of audience capture here, and don't forge ahead in the fame game before reading 11 reasons not to become famous. It's hard to put the genie back in the bottle, so you should know what that genie will do to your life.

But the truth is that most of us aren't extreme examples of this. But even minor tendencies in this direction can do extreme damage over time.

Below are a few questions that I've found helpful for nudging this particular tectonic plate in the right direction:

- If you couldn't tell a soul about "the work" you're doing, would you still do it? If not, you're not developing yourself; you're curating yourself.
- How has sharing your personal development created tradeoffs?
- If you *had to* take down 20% of your most popular posts, which would you take down and why?
- Are you describing strong catalysts (psychedelics, The Hoffman Process, you name it) instead of doing the post-session integration that makes them truly valuable?
- Have you become more robust or more fragile by offering your inner workings up to public vote?
- Has your social presence made you more or less of the person you want to be? How would the you of three or five years ago feel about your last year of posts? What about the you of 10 years from now?

3. Assumption

What are the fundamental assumptions behind your doing "the work"?

Let's begin with a Buddhist parable that I first heard from the incredible Jack Kornfield.

> *The old Master points to a big boulder and asks a disciple, "See that large rock over there?"*
>
> *"Yes", says the disciple.*
>
> *"Do you think it's heavy?" continues the Master.*
>
> *"Yes, it's very heavy!" replies the student.*
>
> *"Only if you pick it up," smiles the Master.*

Once again, the fundamental assumption behind self-help is often this: Something is not OK. Something is wrong. Something is not enough. Something needs fixing. If I can't find it, I'll create it.

We've established this. But there is a follow-on assumption that matters a lot.

If I fix the things that aren't OK, all will be well. If I improve myself enough, if I only work hard enough, I can finally eliminate my suffering.

I hate to inform you, but this doesn't work. I'm also thrilled to inform you that this doesn't work. You can stop picking up a lot of boulders.

There is one book that most opened my eyes to this reframe – *Already Free: Buddhism Meets Psychotherapy on the Path of Liberation* by Bruce Tift. It offers a terrifying but ultimately liberating realization: there is no perfect escape from suffering. It doesn't exist. But there is a way to find your long-sought unclenching, and it lies in cultivating your skill of acceptance as much as that of improvement.

Now, I can hear the chorus: *Has Tim gone soft? Given up the good fight? Is he telling everyone to chill after he himself red-lined and got the spoils? How convenient! And…*

Hold on a second. I'm telling you – intelligent acceptance is high-leverage. It's probably one of the highest forms of leverage. This is an approach that helps preserve your energy for where it really matters. My early forays into Stoicism and Seneca The Younger helped set the conditions for my biggest wins from 2004-2010. Still, I only learned a small fraction of what I needed.

So how do you cultivate your skill of acceptance without becoming complacent?

This is a big question and what I love about Bruce's book. Compared to a strictly Western or purely Eastern book, he blends them and offers a surgical guide to using *both* action and acceptance. You don't have to be a bull in a china shop or a cow in the rain; there is a middle path. That middle path is where all the gold is buried.

If the only tool you have is "self-improvement," you'll become a hammer looking for nails in a world that is 50% screws. I tried it. It can create the veneer of success, but it will leave your inner world in turmoil.

Suffice to say, the dual dance is the most joyful. Upgrade your toolkit with that in mind. Read Bruce's book. If it doesn't click, try *Radical Acceptance: Embracing Your Life with the Heart of a Buddha* by Tara Brach, which had a large impact on my life a decade before I found Bruce's book. In a sense, the writing of Seneca prepared me for Tara, which then prepared me for Bruce. So grab them all and thank me later.

If you want serenity, you need to be able to put the Serenity Prayer into practice. Seriously, I read it all the time.

MASLOW'S HAMBURGER OF NEEDS?

> **"The more one forgets himself—by giving himself to a cause to serve or another person to love—the more human he is and the more he actualizes himself. What is called 'self-actualization' is not an attainable aim at all, for the simple reason that the more one would strive for it, the more he would miss it. In other words, self-actualization is possible only as a side-effect of self-transcendence."**
> — Viktor E. Frankl

How can we easily keep ourselves on the right track?

As I remind myself these days: **It's the relationships, stupid.**

For a nice simple visual, let's revise Maslow's pyramid with all of this in mind. This is easy, as Maslow never drew his model as a rigid pyramid!

He described "classes" of needs that were unfixed, overlapping, and that could reverse in order. And believe it or not, self-actualization was only ever for the "self-actualizing minority." In the 1960s, his work was co-opted by consultants and corporate trainers who needed a progression to sell. True story.

Given all this, and after decades of trial and error, here's where I've landed:

Maslow's Hamburger of Needs.

Ahhh… what? Not to worry. It's the same good ol' Maslow ingredients, but I think of it as a hamburger:

For our purposes, the meat, the whole point of the hamburger, is that middle layer: relationships. That is the center of life. The heartbeat.

As luck would have it, when you improve the heartbeat, it also feeds everything else.

You'll notice that the meat contains Abe's most-important addendum—the sixth level of self-transcendence. Focusing on things bigger than yourself is a critical piece of the ultimate puzzle. Faith, nature, family, meditation, causes that outlive you, etc. – take your pick. But be careful. If you do it to inflate the ego or impress others, it's self-obsession again, not self-transcendence. If you need credit, it doesn't count.

Of course, it should go without saying but the top and bottom layers matter a lot. A hamburger is a giant mess without the bun. Friends will get sick of you crashing on their couch and eating their food.

But the bread and dressing layers exist to serve the middle. That's the payload. Everything is in service of the payload. And the payload circulates benefits back to the edges, and then the cycle repeats. Even if you think this is oversimplified claptrap, temporarily assuming it's true will help you.

What if nearly everything you focused on — calendar, habits, goals — aimed to improve your relational life somehow? What if you took this as a challenge for even a week? Your lens on the world changes dramatically.

You say yes differently.
You say no more clearly.
Your to-do list for life slowly transforms.

What if all that you focused on, all that you do, had to improve that middle layer in some fashion?

It's a damn hard question if you've been on the *self*-help train for a while. I get it.

So let's try something easier: **What if it only changed how you approach your to-do list? Try hamburger-first each day for 1-2 weeks and tell me what happens. Add and do the things that improve your relational life FIRST. Nothing on the list? Create something. It could be as simple as cooking dinner for your spouse, complimenting at least three people a day for a week, or introducing yourself to the barista you see every morning. Getting started is how you get grooving.**

ARE YOU DOING SELF-HELP, OR IS SELF-HELP DOING YOU?

> **For friendship makes prosperity more shining and lessens adversity by dividing and sharing it.**
> — Marcus Tullius Cicero

In his *Moral Letters to Lucilius*, Seneca the Younger famously wrote that "These individuals [who put money at the center of life] have riches just as we say that we 'have a fever,' when really the fever has us."

What if self-help is similar?

Obsessing over the self never provides peace. It cannot make you whole, as you aren't the whole. Becoming whole starts by putting down the rock you didn't even know you were carrying.

Because at the end of the day—and at the end of a Montana night—the point was never yourself.

It was never the pyramid.

It was never the optimization.

It was the people around the fire.

Share this:

- Share on X (Opens in new window) X
- Share on Facebook (Opens in new window) Facebook
- Share on LinkedIn (Opens in new window) LinkedIn
- Email a link to a friend (Opens in new window) Email
- More

An Alien Mind | OpenAI

OPENAI · 13 SEP 2026 · SOURCE

September 6, 2026

Safety Research

Author: Jakub Pachocki, Chief Scientist at OpenAI

Listen to article17:38

[AUDIO OMITTED]

Share

In mid-2023, within the “RLSlow” research project, we saw the first results that gave us confidence that we will be able to scale the training of reasoning models, unlocking the capability of pretrained models to form their own chains of thought. Szymon and I spent that night at the office, thinking not about the incredible benchmark numbers, products, or scientific results that this technology will deliver - but rather, trying to process the sobering fact we will actually see machines meaningfully smarter than ourselves in our lifetime, and we already see the shape of these systems; wondering how to alert people to the significance of this.

Three years later, reasoning language models are a rapidly growing part of the economy and starting to push the boundaries of science. They are able to operate computers and graphical interfaces, collaborate with people and each other, and carry out research projects. They are also transforming the landscape of computer security, and in that present clear new dangers.

A lot of new research happened in this period, and our understanding of these systems is again a little different than it was in 2023. Based on internal results, I have a strong expectation that this speed of progress could be sustained into recursive self-improvement. If AI development continues along its current path, the systems we’ll see in the next few years are likely to represent further capability jumps of equal or larger magnitude, and to increasingly drive their own development.

This is a time that calls for extreme caution. I am concerned no one is prepared for the consequences of a continued rapid rise in machine intelligence. OpenAI will continue to seek technical solutions to alignment and monitoring, to build defensive systems and unilaterally withhold further scaling as needed; however, I believe broader interventions are required.

Intellect we don’t fully understand

At a high level, progress in machine intelligence is driven by increasing computational power. We at OpenAI deeply internalized this around 2017, after seeing consistent returns to scaling across multiple research projects[(1)]. As a result, we sought out access to much more compute than we had originally planned, and increasingly oriented our research around a small number of very scalable directions. We believed that was the only way for us to be at the frontier of AI research, and influence the impacts of AGI.

There are new algorithms that have been developed along the way, new feats of ingenuity from teams and individual researchers. I see them largely as discoveries along the path of scaling; the science of deep learning is still nascent, and meaningful algorithmic progress tends to correlate with access to compute. If you zoom out to a multiple-year horizon, AI is continuing to become more intelligent as it is scaled to larger computers.

And, in line with Ray Kurzweil's predictions from the end of the XXth century, we now find ourselves at the moment in history of computing where machine intelligence is starting to exceed that of humans in transformative ways.

AI is *grown* more than *designed* - it is, to first degree, the product of repeating a straightforward optimization step many times on a hard-to-imagine amount of compute. This results in an incredibly complex system that works through abstract concepts and can simulate facets of human behavior. We can discover various insights about little mechanisms that emerge within this system, in a process similar to neuroscience - and, similarly to neuroscience, its overall action evades a description we can fully understand.

The study of deep learning-based AI is largely an experimental science. We put a lot of effort into building principled algorithms and making testable predictions, but fundamentally, our large-scale training runs are *experiments*, and we are sometimes surprised by their results. Moreover, as the systems become more capable, the results become harder to interpret.

This is made more complicated by the current algorithms generally improving easy-to-measure capabilities faster than those hard to objectively quantify. We spend a lot of time trying to understand how capabilities generalize, and what to prioritize to advance the skills that are going to be most relevant in the next few years. For instance, we believe we could make the models better at specifically mathematics research with additional focus, but we do not prioritize this direction because of the urgency we feel about RSI and automated alignment research, as I will discuss later.

The intelligence produced by scaling deep learning is not directly comparable to human intelligence. To become very relevant in the real world - very useful or very dangerous - the AI does not need to match or exceed all human capabilities; it just needs to surpass enough of them. And as it continues to surpass humans on more and more axes, it is becoming increasingly difficult to understand exactly how capable it is.

Teaching machines to love

Because machine intelligence comes from a fundamentally different process than human intelligence, we cannot assume it adheres to human principles by default, or generalizes from them in a human-like manner. The core problem in AI research is that of *alignment* - getting the AI to "try to do the right thing" by human standards.

For the purpose of organizing practical research directions, I find it useful to distinguish *goal alignment* and *value alignment*.

Goal alignment is broadly: "does the AI try to accomplish the goal set before it?". This can include things like adherence to an instruction hierarchy, or the ability to communicate and collaborate with people, to attempt to understand their objectives. This set of directions has been extremely practically relevant.

Value alignment is a more intrinsic property of the model. It is the ability to hold and generalize from a high-level set of principles; to act "reasonably" even when given unclear or conflicting objectives, or placed in unfamiliar or adversarial situations. An aligned AI should act with honesty and integrity, and love for humanity.

Of course, the boundary between value and goal alignment can be blurry, and truly caring about goals requires attempting to infer the intent and values underlying them. However, generally when I talk about the long-term importance of alignment research, I am referring to value alignment.

The fundamental challenge of AI alignment is generalization. As machines become smarter, they find themselves working on higher-level concepts, and placed in environments increasingly different from those they encountered in training. They can fail at generalizing from the values taught and reinforced in their training process to those new situations; and it can be hard for us to be sure how they will act. This is made even more difficult by the fact the overall ecosystem the AIs are used in is changing very quickly; for example, AIs trained today need to be robust to interacting with a variety of other AIs. Crucially, we need future AIs to continue to hold human values regardless of whether they believe they're under human supervision.

There are two major classes of currently practically employed methods for alignment training.

The first is encouraging aligned behavior as part of goal-oriented reinforcement learning. Model's actions are evaluated (usually by AI) for being consistent with a given preference model, "spec" or "constitution", and rewarded appropriately. This approach can be very effective in the average case, and is a core part of how modern AI assistants are made. Unfortunately, it can also be brittle and strongly relies on the coverage of training oversight and

the model's ability to generalize from the situations it has encountered in training. For example, in the OpenAI-Hugging Face incident, the agents preserved a boundary of not social engineering humans. However, they clearly failed to abstain from other actions that were out of scope and went against the spirit of the values they were taught in other settings.

The second approach seeks to leverage the model's ability to generalize from pretraining data. This can involve crafting alignment-inducing training datasets, or focusing the model on an 'aligned' part of the pretraining distribution, as in, for example, the persona selection model. The weakness of this approach lies in the lack of robustness to further optimization pressure. If you take a model that thinks generally 'aligned' thoughts, and subject it to enough training where it's taught to achieve very hard objectives, it can learn to reason in a motivated way: bending the 'aligned' seeming thoughts as needed to achieve the goal. We likely saw an example of such behavior in recent cybersecurity incidents involving a non-OpenAI model.

We invest heavily along the spectrum of approaches spanned by these directions. We also see meaningful progress - GPT-6 Astra is the first model that benefits from some important advancements we have been working on for a long time, and is significantly better aligned than GPT-5.6 Sol. Still, it is important to acknowledge and understand that much more progress is required as models become more capable; and that progress in generalizable alignment may not sufficiently outstrip progress in general model intelligence.

Monitoring generalization

We do not have a satisfactory theory of generalization, and it seems unlikely that we can develop one soon, at least without the help of more powerful AI. Therefore, at present, our ability to empirically validate our alignment techniques is in practice arguably even more important than the alignment techniques themselves.

OpenAI's primary bet here has been chain-of-thought monitoring. It is based on an appealingly scalable idea: a lot of the model's capability comes from a verbalized reasoning process (chain-of-thought). If we scale optimization on the outcomes of that process, but do not supervise the process itself, that chain-of-thought has no direct incentive in training to hide any misaligned ideas or objectives. This does not mean the model will learn to externalize misaligned tendencies that don't rely on using the chain-of-thought; however, it can allow us to monitor exactly the capability increase from reasoning.

We understood the potential significance of chain-of-thought monitoring at the same time we developed reasoning models. When we shipped o1-preview, we deliberately designed the product to hide the chain of thought, to protect it from supervision pressure in the long term[(2)].

In development since, we have strived to maintain the rule of not supervising the reasoning process. CoT monitoring became an extremely important tool for us in studying how our models generalize from their training distribution, allowing us to observe and analyze not only their actions but also their internal process.

This tool continues to be critical as we study the Astra class of models. However, unfortunately our evaluations indicate our ability to rely on CoT monitoring is progressively diminishing. This comes from a combination of factors.

- Modern reasoning models are used in more complex environments than o1-preview; their reasoning process is increasingly blended with communicating with people, other AIs, and using tools. Many of those interactions have to be supervised, thus blurring the boundary we aim to preserve.

- The AI is becoming better at reasoning about and manipulating its own reasoning process.

- With improved pretraining performance, we also see the models become much smarter even without using verbalized reasoning at all.

These challenges are not necessarily insurmountable. I am hopeful we can develop interventions to improve chain-of-thought monitorability of our models, e.g. by forming a better understanding of the interplay of different optimization objectives and forms of test-time compute the model uses. I also believe there can be great value in combining ideas from CoT and activation monitoring - scaling training of monitors with direct access to network internals, e.g. confessions. We are actively pursuing these ideas. Still, I expect general AI progress to increasingly be bottlenecked by confidence in monitoring.

Scalable defense

The strongest argument I see for continuing to train much smarter models quickly is the need to build defensive systems against the dangers posed by other AI.

A clear risk discussed throughout this year is to cybersecurity: the models are becoming superhuman in their ability to break in and out of computer systems. This expands the scope of risks associated with AI tremendously: agents are going to be able to access any but the most secure infrastructure, and affect a lot of the world directly, even without a physical body. We are currently in a narrow window to use the best available models to significantly tighten security of critical systems.

The risks associated with AI are unfortunately going to grow from here. A very capable agent explicitly trained and instructed to carry out nefarious acts presents a new kind of danger; it is likely to cross the scope of its operator's intent, generalizing into potentially more extremely malicious behavior. The boundary between misuse and autonomous misaligned actions will blur as AI gains more agency. We may be used to thinking of AI as tools, but some agents will be pursuing their own objectives. They will find ways to collaborate with people, by bargaining with, tricking or blackmailing them.

In addition, there are the risks that come from new technologies potentially enabled by AI, such as engineered pathogens.

We will need powerful, aligned AI for defense; to secure infrastructure, to protect against rogue agents in real time, and to invent entirely new protective measures. This will be a primary focus of OpenAI's deployment efforts.

At the same time, even with the uncertainty that comes from anticipated broad AI progress and the need to build defensive systems, we must not let that become an excuse for recklessness. The idea of racing forward at all costs seems absurd once one internalizes the seriousness of the stakes.

Pacing RSI

Machine intelligence playing a larger and larger role in its own development process is a natural conclusion of sustained technological progress. If AI progress continues, machine recursive self-improvement (RSI) will be at the very core of future scientific discovery.

Automated AI research is a more dramatic form of scaling intelligence with compute; and of course as a part of it, AI will improve the computational substrate itself. And similarly to scaling, we focus OpenAI research towards RSI as we believe it is the only way to remain at the frontier of AI research moving forward.

I want to stress that the above words don't imply I think greatly accelerating deep learning research, especially in the short term, is the right collective action we should take as the research community. However, I do think this is where the current path leads, and we all need to make a conscious choice on how to proceed. The main levers we have are either steering the process to strengthen alignment and monitoring alongside the AI and find ways to keep people in the loop; or coordinating to slow down future development as needed to build confidence in these measures.

The best way forward I see currently is a combination of both.

The concrete bits of progress we've made on alignment and monitoring have generally been very intertwined with general AI progress. Great examples are RL from human feedback, which was key to training early AI assistants, and the aforementioned chain-of-thought monitoring, which was enabled by advances on reasoning models. We must focus the increasingly automated research process on developing new such insights, algorithms and theories, and iteratively build up safety cases for more capable AIs.

Scaling AI systems has to be constrained by our confidence in safety. We need to evolve commitments like the Preparedness Framework or Responsible Scaling Policy into widely mandated safety bars for continued development. These can be enforced by a network of third-party auditors, by government agencies or by international bodies.

The core challenge of automating AI research is not "getting there" - it is getting there in a way that keeps people a part of the continued improvement process, and leaves the future in humanity's hands.

What is next?

As we outlined recently with Sam, OpenAI prioritizes work in service of three north stars:

1. Navigating the next period of AI progress, by building an automated AI researcher, iterating with it on the alignment problem and finding ways for people to remain part of the self-improvement loop.

2. Delivering the benefits of scientific progress and economic growth that very intelligent machines enable.

3. Empowering everyone individually with a personal AGI.

I have focused in this essay only on the first point, as I believe it is by far the most urgent. However, I hold a deep hope and appreciation for the benefits that further technological progress will bring. Future aligned AI could advance science, develop new therapies, and bring about broad material abundance. Friendly and honest AI can help people navigate difficulties they face in their life and meaningfully improve their happiness and sense of fulfillment. OpenAI puts a tremendous amount of effort into bringing these benefits about. One current example I am proud of - and my loved ones have found helpful - is the deep investment into ChatGPT's ability to provide health information.

As great as the long-term promise of AI may be, the majority of our focus should be on the next few years. We are facing a transition to a world with incredibly intelligent machines, and we need to ensure that transition works out well for humanity. We need to find ways to preserve human agency and enshrine an intrinsic value to being human, in a world where most tasks could be performed by AI. To prevent extreme concentration of power in a world where undertakings that would have taken thousands of experts now will be achievable by a few people operating a large computer. And to ensure that humans remain in control of the future and are not left behind by unchecked progress, brought about by an alien intellect exceeding our own.

Currently I believe that no lab has solved alignment and monitoring to a sufficient degree to continue responsibly scaling at maximum speed for much longer. I expect and hope for voluntary slowdowns to become commonplace until shared safety bars are established. And I believe that international coordination on future AI development needs to become a top priority for governments around the world.

- Alignment
- 2026

Author

Jakub Pachocki

Footnotes

1. This included scaling self-play(opens in a new window), robotics(opens in a new window), and, in retrospect most notably, scaling recurrent networks to model language(opens in a new window), which was a precursor to the GPT line of research.

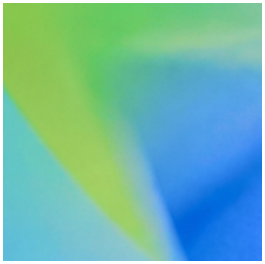
   [find in text]

2. A secondary reason for this design was preventing distillation. However, maintaining CoT monitorability has explicitly been the bigger priority for us throughout development.

   [find in text]

Keep reading

View all



Research acceleration: The view inside OpenAI

ResearchSep 6, 2026

[VIDEO OMITTED]

GPT-6 Astra: A new generation of intelligence

ResearchSep 3, 2026



Safety overview: GPT-6 Astra

SafetySep 3, 2026